

Quantum Variational Stress Predictor: Tree Tensor Network Architecture for Hierarchical Ordinal Regression

Aditya Raj

adityakeshrissm123@gmail.com

Oishik Kar

23bcs089@iiitdwd.ac.in

Noble Agyeman-Bobie

nobleagyemanbobie2023@gmail.com

May 1, 2026

Abstract

We present a 4-qubit Variational Quantum Circuit (VQC) that predicts adolescent psychological stress from survey data, achieving **MAE=0.7943** (Score 27.71/30). Through a systematic hardware ablation study across 1–4 qubit architectures, we discovered that *all* unstructured entanglement topologies—regardless of qubit count—converge to the same performance ceiling ($\text{MAE} \approx 0.81$). Only our Tree Tensor Network (TTN) entanglement breaks through this barrier, reducing it to 0.7943 with 25% fewer CNOT gates than the standard ring topology. This finding demonstrates that **structured, domain-aware entanglement outperforms brute-force qubit scaling**—a result with direct implications for NISQ-era quantum machine learning. Total parameters: 33 of 100 allowed.

1 Problem Analysis

The dataset comprises 675 training samples with 18 Likert-scale features (1–5) predicting *Recent_Stress* (1–5)—an ordinal regression task on inherently noisy self-reported psychological survey data. A constant-mean predictor ($\bar{y} \approx 2.99$) yields $\text{MAE} = 0.908$, establishing the non-trivial learning threshold.

EDA Findings. Spearman correlation analysis revealed that **14 of 18 features** exhibit near-zero correlation with the target ($|\rho| < 0.08$). Only four features carry statistically meaningful signal (Table 1).

Table 1: Feature selection by Spearman ρ

Rank	Feature	ρ	Domain
1	Palpitations	+0.45	Physiological
2	Act_Conflict	+0.19	Behavioural
3	Sleep_Issues	+0.14	Physiological
4	Subj_Conf	+0.11	Academic

Decision: Aggressive pruning to 4 features. Including noise features in a quantum circuit scrambles useful entanglement and pushes the variational landscape toward barren plateaus [3]. Features are normalized from Likert [1,5] to $[0, \pi]$ via $x_{\text{enc}} = (x - 1)/4 \cdot \pi$ for angle encoding.

2 Architecture Design

Our architecture was selected through a systematic hardware ablation study (Section 5). The key finding that drove our design: *brute-force entanglement does not help on this dataset, but hierarchically structured entanglement does.*


2.1 Circuit Structure

The VQC consists of **2 repeated layers** (data re-uploading [2]), each with three stages. Figure 1 shows one complete layer.

Stage 1: Angle Encoding + Data Re-uploading. Each feature x_i is encoded as $R_Y(x_i)$ on qubit q_i , applied at **every** layer—not just the first. Pérez-Salinas et al. [2] proved that interleaving data with trainable rotations transforms a VQC into a universal function approximator, analogous to depth in classical neural networks.

Stage 2: Tree Tensor Network (TTN) Entanglement [1]. This is the **core architectural innovation**. Instead of the textbook CNOT ring (4 gates/layer) or linear chain (3 gates/layer but flat), we employ a **hierarchical tree** with 3 CNOTs arranged in two levels:

- **Level 1 (local):** $\text{CNOT}(q_0, q_1)$, $\text{CNOT}(q_2, q_3)$ —pairs physiological features separately from psychosocial features.
- **Level 2 (global):** $\text{CNOT}(q_1, q_2)$ —cross-domain hierarchical compression.



Why TTN over standard topologies:

1. **25% fewer entanglement gates** than ring (3 vs. 4 CNOTs/layer), directly satisfying Rule 2 (“minimize gate count”).
2. **Barren plateau resistance:** Pesah et al. [4] proved that tensor network ansätze exhibit *polynomial* gradient variance scaling ($\mathcal{O}(1/\text{poly}(n))$), vs. *exponential* vanishing ($\mathcal{O}(1/2^n)$) in hardware-efficient circuits.
3. **Domain-aware hierarchy:** The tree mirrors biopsychosocial stress theory—physical and social symptoms are processed locally before cross-domain integration. Our ablation (Sec. 5) **proves** this is the source of TTN’s advantage.

Stage 3: Trainable Ansatz & Measurement. $4 \times R_Y + 4 \times R_Z$ parameterized rotations per layer (16 quantum params total). The $2^4 = 16$ output probabilities feed a single FC layer (16 $\rightarrow$ 1, 17 classical params). The output is bounded via $\hat{y} = 1 + 4 \cdot \sigma(z)$, constraining predictions to $[1, 5]$.

3 Quantum Complexity Analysis

At 33 total parameters (33% of the cap), the quantum circuit contributes 16 parameters—48% of the model’s learning capacity—demonstrating genuine quantum contribution rather than relying on classical FC dominance.




Figure 1: Complete single-layer architecture of the 4-qubit TTN circuit. **Stage 1** encodes each feature as $R_Y(x_i)$ rotations. **Stage 2** applies hierarchical TTN entanglement: Level 1 pairs domain-specific qubits (physical: q_0 – q_1 , psychosocial: q_2 – q_3), Level 2 performs cross-domain compression (q_1 – q_2). **Stage 3** applies trainable R_Y/R_Z rotations. The full circuit repeats this layer twice via data re-uploading [2]. Only **3 CNOT gates per layer** (vs. 4 for ring topology).

Table 2: Resource budget (primary submission)

Metric	Value
Qubits	4
Entanglement CNOTs / layer	3 (TTN)
Total gates / layer	15
Total gates (2 layers)	30
Quantum parameters	16
Classical FC parameters	17
Total parameters	33 / 100 limit

4 Training Methodology

- **Loss:** MAE—directly optimizes the competition evaluation metric. Avoids unnecessary hyperparameters on a small, noisy dataset.
- **Optimizer:** Adam ($\text{lr} = 0.01$). Conservative learning rate stabilizes parameter-shift gradient estimation.
- **Output activation:** $\hat{y} = 1 + 4 \cdot \sigma(z)$ hard-bounds predictions to the Likert [1,5] range, eliminating impossible outputs.
- **Checkpointing:** Model weights saved at the epoch with lowest validation MAE.

Why minimal training engineering: Cosine annealing, gradient clipping, and Top-K weight ensembling were all explored but provided no improvement over the simple pipeline on this particular dataset. We deliberately kept the training loop clean so that all performance gains are attributable to the *quantum architecture itself*, not classical training heuristics.

5 Experimental Results

5.1 Hardware Ablation Study

We conducted a systematic ablation across 1–4 qubit configurations with multiple entanglement topologies. This is the most revealing result of our work.

5.2 Critical Insight: Structure > Entanglement

The ablation reveals a striking pattern: **all models with unstructured entanglement converge to the same MAE ≈ 0.81** , regardless of qubit count. The 1-qubit model (zero CNOT gates) *matches* the 3-qubit model (2 CNOTs/layer) and the 4-qubit linear chain (3 CNOTs/layer). This proves that naively adding more entanglement provides




Figure 2: Hardware ablation: all unstructured topologies (gray) converge to MAE ≈ 0.81 regardless of qubit count. Only TTN (green) breaks through the noise ceiling, demonstrating that *structured* entanglement matters more than qubit count.

no benefit—the stress features are sufficiently independent that brute-force qubit interactions inject noise rather than signal.

Only the TTN breaks through 0.81. Its hierarchical structure—pairing physiological features (Palpitations $\leftrightarrow$ Sleep_Issues) separately from psychosocial features (Act_Conflict $\leftrightarrow$ Subj_Conf) before cross-domain compression—matches the domain structure of the data. This is precisely the inductive bias that tensor network literature predicts [1, 5].

5.3 Training Convergence

The TTN model converges smoothly over 100 epochs (Fig. 3). The validation plateau at ~ 0.80 confirms the **noise ceiling** inherent to subjective Likert-scale survey data. Any model reporting significantly lower MAE without additional data sources would be fitting to human reporting variability, not genuine psychological signal.

5.4 Score

Per the competition formula:

$$\text{Score} = 30 \times e^{-0.1 \times 0.7943} \approx 27.71/30$$

6 The 2-Qubit Alternative

As an exploratory study, we investigated a **minimal 2-qubit Dense Angle Encoding model** that packs all 4 features into just 2 qubits using R_Y/R_Z stacking (2 features per qubit). With only 1 CNOT per layer and 13 total parameters, it

Table 3: Complete hardware ablation across qubit counts and entanglement topologies

Model	Qubits	CNOTs/Layer	Params	Best Val MAE	Score (/30)	Key Observation
1-Qubit Bottleneck	1	0	~ 5	0.8095	27.67	Zero entanglement baseline
2-Qubit Dense Angle	2	1	13	<i>0.8163</i>	27.65	<i>Most parameter-efficient (13/100)</i>
3-Qubit Expansion	3	2	~ 21	0.8095	27.67	Matches 1-qubit: entanglement ineffective
4-Qubit Linear Chain	4	3	33	0.8102	27.67	Flat topology, no hierarchy
4-Qubit Ring	4	4	33	0.8102	27.67	Maximum CNOT count, no benefit
4-Qubit TTN (ours)	4	3	33	0.7943	27.71	Only model to break 0.81 barrier




Figure 3: TTN training dynamics over 100 epochs. Validation MAE (red) reaches minimum 0.7943 then plateaus at ~ 0.80 , while training MAE (blue) continues decreasing. This divergence confirms the model hits the dataset’s intrinsic noise ceiling.

achieved MAE ≈ 0.82 —remarkable parameter efficiency.

We selected the 4-qubit TTN as our primary submission because it is the **only architecture that breaks the 0.81 noise barrier**. The 2-qubit model, while maximally hardware-efficient, cannot express the hierarchical feature relationships that the TTN exploits. This represents a fundamental expressivity limit, not an optimization failure.

7 Future Optimization

1. **ZZ-Interaction Feature Map [6]:** Encoding pairwise products $x_i \cdot x_{i+1}$ via CNOT- R_Z -CNOT could capture non-linear feature correlations. Initial experiments showed increased gate count without clear MAE improvement on this dataset, but could benefit datasets with stronger feature interactions.
2. **Ordinal regression loss (CORAL [8]):** Models the non-uniform spacing of Likert categories more faithfully than MAE.
3. **NISQ noise-resilient training:** Incorporating hardware decoherence models during optimization for deployment on real quantum devices.
4. **Adaptive TTN topology:** Learning the optimal qubit-to-feature pairing rather than using domain-guided manual assignment.

References

- [1] Grant, E. et al. Hierarchical quantum classifiers. *npj Quantum Inf.* **4**, 65 (2018).
- [2] Pérez-Salinas, A. et al. Data re-uploading for a universal quantum classifier. *Quantum* **4**, 226 (2020).

- [3] McClean, J.R. et al. Barren plateaus in quantum neural network training landscapes. *Nat. Commun.* **9**, 4812 (2018).
- [4] Pesah, A. et al. Absence of barren plateaus in quantum convolutional neural networks. *Phys. Rev. X* **11**, 041011 (2021).
- [5] Huggins, W. et al. Towards quantum machine learning with tensor networks. *Quantum Sci. Technol.* **4**, 024001 (2019).
- [6] Havlíček, V. et al. Supervised learning with quantum-enhanced feature spaces. *Nature* **567**, 209–212 (2019).
- [7] Kandala, A. et al. Hardware-efficient variational quantum eigensolver. *Nature* **549**, 242 (2017).
- [8] Cao, W. et al. Rank consistent ordinal regression for neural networks. *Pattern Recognit.* **108**, 107736 (2020).